

# ST 558 HW 5

Alise Miller

## Task 1: Conceptual Questions

1. The purposes of cross-validation when fitting a random forest model is to help find the best fit. So find the model that has the best metrics and that doesn't over fit for best generalization.
2. Bagged tree algorithm has many steps. The steps include having your original sample. From there getting the bootstrap samples. For non parametric we treat the sample as population, the resampling is done with replacement and we can get the same observation multiple times. After all of the trees are trained we get the bootstrap statistics.
3. General linear model is  $y = B_0 + B_1X_1 + B_2X_2 + B_3X_3 \dots$ . This means a model that uses one or more predictors to predict the response variable. When I think of GLM is the the first model that comes to mind when making a model.
4. Adding an interaction term to MLR helps control the effect of one predictor to another. Predictors can change based on the level of other predictors.
5. We split our data into training and test sets so that we can have an accurate model and so that the data isn't over fit. The training set helps to make the model, while the testing evaluate how well the model will work for future data.

## Task 2: Data Prep

### packages and data

```
suppressWarnings(library(tidyverse))
suppressWarnings(library(tidymodels))
suppressWarnings(library(caret))
suppressWarnings(library(yardstick))

heart_data <- read_csv("https://www4.stat.ncsu.edu/~online/datasets/heart.csv")
```

## 1. Run and report summary

```
summary(heart_data)
```

Age	Sex	ChestPainType	RestingBP
Min. :28.00	Length:918	Length:918	Min. : 0.0
1st Qu.:47.00	Class :character	Class :character	1st Qu.:120.0
Median :54.00	Mode :character	Mode :character	Median :130.0
Mean :53.51			Mean :132.4
3rd Qu.:60.00			3rd Qu.:140.0
Max. :77.00			Max. :200.0
Cholesterol	FastingBS	RestingECG	MaxHR
Min. : 0.0	Min. :0.0000	Length:918	Min. : 60.0
1st Qu.:173.2	1st Qu.:0.0000	Class :character	1st Qu.:120.0
Median :223.0	Median :0.0000	Mode :character	Median :138.0
Mean :198.8	Mean :0.2331		Mean :136.8
3rd Qu.:267.0	3rd Qu.:0.0000		3rd Qu.:156.0
Max. :603.0	Max. :1.0000		Max. :202.0
ExerciseAngina	Oldpeak	ST_Slope	HeartDisease
Length:918	Min. :-2.6000	Length:918	Min. :0.0000
Class :character	1st Qu.: 0.0000	Class :character	1st Qu.:0.0000
Mode :character	Median : 0.6000	Mode :character	Median :1.0000
	Mean : 0.8874		Mean :0.5534
	3rd Qu.: 1.5000		3rd Qu.:1.0000
	Max. : 6.2000		Max. :1.0000

```
str(heart_data)
```

```
spc_tbl_ [918 x 12] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
 $ Age      : num [1:918] 40 49 37 48 54 39 45 54 37 48 ...
 $ Sex      : chr [1:918] "M" "F" "M" "F" ...
 $ ChestPainType : chr [1:918] "ATA" "NAP" "ATA" "ASY" ...
 $ RestingBP  : num [1:918] 140 160 130 138 150 120 130 110 140 120 ...
 $ Cholesterol : num [1:918] 289 180 283 214 195 339 237 208 207 284 ...
 $ FastingBS  : num [1:918] 0 0 0 0 0 0 0 0 0 0 ...
 $ RestingECG : chr [1:918] "Normal" "Normal" "ST" "Normal" ...
 $ MaxHR      : num [1:918] 172 156 98 108 122 170 170 142 130 120 ...
 $ ExerciseAngina: chr [1:918] "N" "N" "N" "Y" ...
 $ Oldpeak    : num [1:918] 0 1 0 1.5 0 0 0 0 1.5 0 ...
 $ ST_Slope   : chr [1:918] "Up" "Flat" "Up" "Flat" ...
```

```

$ HeartDisease : num [1:918] 0 1 0 1 0 0 0 0 1 0 ...
- attr(*, "spec")=
.. cols(
..   Age = col_double(),
..   Sex = col_character(),
..   ChestPainType = col_character(),
..   RestingBP = col_double(),
..   Cholesterol = col_double(),
..   FastingBS = col_double(),
..   RestingECG = col_character(),
..   MaxHR = col_double(),
..   ExerciseAngina = col_character(),
..   Oldpeak = col_double(),
..   ST_Slope = col_character(),
..   HeartDisease = col_double()
.. )
- attr(*, "problems")=<externalptr>

```

1. a. In R Heart Disease is “num” so it is understood as quantitative. 1.b. This doesn’t make sense because these are it’s yes or no and it does not make sense to average yes or no. Maybe count or percent of Yes versus No’s makes better sense.

## 2. Change Heart Disease

```

new_heart <- heart_data |>
  mutate(heartdisease_yes_no =
    factor(HeartDisease,
      levels = c(0,1),
      labels = c("No", "Yes"))
  ) |>
  select(-HeartDisease, -ST_Slope)
glimpse(new_heart) #just to check

```

Rows: 918

Columns: 11

```

$ Age      <dbl> 40, 49, 37, 48, 54, 39, 45, 54, 37, 48, 37, 58, 39~
$ Sex      <chr> "M", "F", "M", "F", "M", "M", "F", "M", "M", "F", ~
$ ChestPainType <chr> "ATA", "NAP", "ATA", "ASY", "NAP", "NAP", "ATA", "~
$ RestingBP <dbl> 140, 160, 130, 138, 150, 120, 130, 110, 140, 120, ~
$ Cholesterol <dbl> 289, 180, 283, 214, 195, 339, 237, 208, 207, 284, ~
$ FastingBS <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ~

```

```

$ RestingECG      <chr> "Normal", "Normal", "ST", "Normal", "Normal", "Nor~
$ MaxHR           <dbl> 172, 156, 98, 108, 122, 170, 170, 142, 130, 120, 1~
$ ExerciseAngina  <chr> "N", "N", "N", "Y", "N", "N", "N", "N", "Y", "N", ~
$ Oldpeak         <dbl> 0.0, 1.0, 0.0, 1.5, 0.0, 0.0, 0.0, 0.0, 1.5, 0.0, ~
$ heartdisease_yes_no <fct> No, Yes, No, Yes, No, No, No, No, Yes, No, Yes~

```

## Task 3 EDA


### 1. Plot

```

ggplot(new_heart,
      aes(x= MaxHR, y = Age, color = heartdisease_yes_no)) +
  geom_point() +
  geom_smooth(method = "lm", se = FALSE) +
  labs(
    x= "Max Heart Rate",
    y= "Age",
    title = "Age Predicted by Max HR and Heart Disease",
    color = "Heart Disease") +
  scale_color_viridis_d()

```

```
`geom_smooth()` using formula = 'y ~ x'
```



## 2. Interpret Plot

It seems like both lines have a downward slope, but they intersect so they are not the same slope. From the scatter plot it seems like people that have heart disease are older and have a lower max heart rate. While people with no heart disease are younger and have a higher max heart rate. The line for people with and without heart disease seem to have a similar y-intercept with the no heart disease being slightly higher. With all of this I think that it is best we use an interaction model.

## Task 4 Testing and Training

```
set.seed(101)
new_heart_split <- initial_split(new_heart, prop = 0.8)
new_heart_train <- training (new_heart_split)
new_heart_test <- testing (new_heart_split)
```

## Task 5: OLS and LASSO

### 1. Interaction Model

```
ols_mlr <- lm(formula = Age ~ MaxHR + heartdisease_yes_no
              + MaxHR*heartdisease_yes_no,
              data = new_heart_train )
summary(ols_mlr)
```

Call:

```
lm(formula = Age ~ MaxHR + heartdisease_yes_no + MaxHR * heartdisease_yes_no,
    data = new_heart_train)
```

Residuals:

Min	1Q	Median	3Q	Max
-22.7703	-5.7966	0.4516	5.7772	20.6378

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	75.58896	3.07510	24.581	< 2e-16 ***
MaxHR	-0.16992	0.02064	-8.233	8.43e-16 ***
heartdisease_yes_noYes	-8.58502	3.83433	-2.239	0.02546 *
MaxHR:heartdisease_yes_noYes	0.08343	0.02716	3.072	0.00221 **

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 8.478 on 730 degrees of freedom

Multiple R-squared: 0.1839, Adjusted R-squared: 0.1806

F-statistic: 54.84 on 3 and 730 DF, p-value: < 2.2e-16

### 2. Predict and RMSE

```
pred1 <- predict(ols_mlr, newdata = new_heart_test)
results1 <- data.frame(.pred = pred1, Age = new_heart_test$Age)

OLS_rmse <- rmse(results1, truth = Age, estimate = .pred)
```

### 3. LASSO Recipe

```
new_heart_cv_folds <- vfold_cv(new_heart_train, 10)
LASSO_recipe <-
  recipe( Age ~ MaxHR + heartdisease_yes_no, data = new_heart_train) |>
  step_dummy(all_nominal_predictors()) |>
  step_normalize(all_numeric_predictors()) |>
  step_interact(terms = ~ MaxHR:starts_with("heartdisease_yes_no"))
LASSO_recipe
```

-- Recipe -----

-- Inputs

Number of variables by role

outcome: 1  
predictor: 2

-- Operations

\* Dummy variables from: all\_nominal\_predictors()

\* Centering and scaling for: all\_numeric\_predictors()

\* Interactions with: MaxHR:starts\_with("heartdisease\_yes\_no")

### 4. LASSO Spec and Grid

```
LASSO_spec <- linear_reg(penalty = tune(), mixture = 1) |>
  set_engine("glmnet")

LASSO_wkf <- workflow() |>
  add_recipe(LASSO_recipe) |>
  add_model(LASSO_spec)
LASSO_wkf
```

```
== Workflow =====
Preprocessor: Recipe
Model: linear_reg()

-- Preprocessor -----
3 Recipe Steps

* step_dummy()
* step_normalize()
* step_interact()

-- Model -----
Linear Regression Model Specification (regression)

Main Arguments:
  penalty = tune()
  mixture = 1

Computational engine: glmnet
```

```
LASSO_grid <- LASSO_wkf |>
  tune_grid(resamples = new_heart_cv_folds,
            grid = grid_regular(penalty(), levels = 200),
            metrics = metric_set(rmse))
```

Warning: package 'glmnet' was built under R version 4.5.1

```
LASSO_grid
```

```
# Tuning results
# 10-fold cross-validation
# A tibble: 10 x 4
```



	splits <list>	id <chr>	.metrics <list>	.notes <list>
1	<split [660/74]>	Fold01	<tibble [200 x 5]>	<tibble [0 x 3]>
2	<split [660/74]>	Fold02	<tibble [200 x 5]>	<tibble [0 x 3]>
3	<split [660/74]>	Fold03	<tibble [200 x 5]>	<tibble [0 x 3]>
4	<split [660/74]>	Fold04	<tibble [200 x 5]>	<tibble [0 x 3]>
5	<split [661/73]>	Fold05	<tibble [200 x 5]>	<tibble [0 x 3]>
6	<split [661/73]>	Fold06	<tibble [200 x 5]>	<tibble [0 x 3]>
7	<split [661/73]>	Fold07	<tibble [200 x 5]>	<tibble [0 x 3]>
8	<split [661/73]>	Fold08	<tibble [200 x 5]>	<tibble [0 x 3]>
9	<split [661/73]>	Fold09	<tibble [200 x 5]>	<tibble [0 x 3]>
10	<split [661/73]>	Fold10	<tibble [200 x 5]>	<tibble [0 x 3]>

```
LASSO_grid |>
  collect_metrics() |>
  filter(.metric == "rmse")
```

```
# A tibble: 200 x 7
  penalty .metric .estimator mean      n std_err .config
  <dbl> <chr>    <chr>    <dbl> <int>  <dbl> <chr>
1 1e-10 rmse      standard  8.50    10  0.162 Preprocessor1_Model001
2 1.12e-10 rmse      standard  8.50    10  0.162 Preprocessor1_Model002
3 1.26e-10 rmse      standard  8.50    10  0.162 Preprocessor1_Model003
4 1.41e-10 rmse      standard  8.50    10  0.162 Preprocessor1_Model004
5 1.59e-10 rmse      standard  8.50    10  0.162 Preprocessor1_Model005
6 1.78e-10 rmse      standard  8.50    10  0.162 Preprocessor1_Model006
7 2.00e-10 rmse      standard  8.50    10  0.162 Preprocessor1_Model007
8 2.25e-10 rmse      standard  8.50    10  0.162 Preprocessor1_Model008
9 2.52e-10 rmse      standard  8.50    10  0.162 Preprocessor1_Model009
10 2.83e-10 rmse      standard  8.50    10  0.162 Preprocessor1_Model010
# i 190 more rows
```

```
lowest_rmse <- LASSO_grid |>
  select_best(metric = "rmse")

lowest_rmse
```

```
# A tibble: 1 x 2
  penalty .config
  <dbl> <chr>
1 0.0174 Preprocessor1_Model165
```

```
LASSO_wkf |>
  finalize_workflow(lowest_rmse)
```

```
== Workflow =====
Preprocessor: Recipe
Model: linear_reg()

-- Preprocessor -----
3 Recipe Steps

* step_dummy()
* step_normalize()
* step_interact()

-- Model -----
Linear Regression Model Specification (regression)

Main Arguments:
  penalty = 0.0174263338600965
  mixture = 1

Computational engine: glmnet
```

```
LASSO_final <- LASSO_wkf |>
  finalize_workflow(lowest_rmse) |>
  fit(new_heart_train)

LASSO_rmse_val <- LASSO_final |>
  predict(new_data = new_heart_test) |>
  pull(.pred) |>
  yardstick::rmse_vec(truth = new_heart_test$Age)

LASSO_rmse_val
```

```
[1] 9.095981
```

```
tidy(LASSO_final)
```

```
# A tibble: 4 x 3
  term                estimate penalty
```

	<chr>	<dbl>	<dbl>
1	(Intercept)	54.0	0.0174
2	MaxHR	-3.08	0.0174
3	heartdisease_yes_no_Yes	1.36	0.0174
4	MaxHR_x_heartdisease_yes_no_Yes	1.03	0.0174

## 5. RMSE

I think that the RMSE calculation to be different, because I think the LASSO would be more accurate. IT has more “training” than the OLS.

## 6. RMSE OLS vs LASSO

```
OLS_rmse
```

```
# A tibble: 1 x 3
  .metric .estimator .estimate
  <chr>    <chr>         <dbl>
1 rmse    standard         9.10
```

```
LASSO_rmse_val
```

```
[1] 9.095981
```

The RMSE for OLS is 9.10 while the RMSE for the LASSO data is 9.096. Although they are LASSO is slightly better, they are roughly the same if you keep 2 decimal places.

## 7. Why are they similar?

I think they are similar because the penalty of the LASSO was very small. So that might mean that there wasn’t anything to shrink.

## Task 6: Logistic Regression

```

LR1_recipe <- recipe(heartdisease_yes_no ~ Sex + Age, data = new_heart_train) |>
  step_normalize(Age) |>
  step_dummy(Sex)

LR2_recipe <- recipe(heartdisease_yes_no ~ Sex + ChestPainType + Age, data = new_heart_train) |>
  step_normalize(all_numeric_predictors()) |>
  step_dummy(all_nominal_predictors())
LR_spec <- logistic_reg() |>
  set_engine("glm")
LR1_wf <- workflow() |>
  add_recipe(LR1_recipe) |>
  add_model(LR_spec)
LR2_wf <- workflow() |>
  add_recipe(LR2_recipe) |>
  add_model(LR_spec)
LR1_fit <- fit(LR1_wf, data = new_heart_test)
LR2_fit <- fit(LR2_wf, data = new_heart_test)

LR1_fit

```

```

== Workflow [trained] =====
Preprocessor: Recipe
Model: logistic_reg()

-- Preprocessor -----
2 Recipe Steps

* step_normalize()
* step_dummy()

-- Model -----

Call:  stats::glm(formula = ..y ~ ., family = stats::binomial, data = data)

Coefficients:
(Intercept)      Age      Sex_M
    -1.3794     0.5833     1.7050

Degrees of Freedom: 183 Total (i.e. Null);  181 Residual
Null Deviance:      255
Residual Deviance: 222.4    AIC: 228.4

```

## LR2\_fit

```
== Workflow [trained] =====  
Preprocessor: Recipe  
Model: logistic_reg()
```

```
-- Preprocessor -----  
2 Recipe Steps
```

```
* step_normalize()  
* step_dummy()
```

```
-- Model -----
```

```
Call: stats::glm(formula = ..y ~ ., family = stats::binomial, data = data)
```

Coefficients:

(Intercept)	Age	Sex_M	ChestPainType_ATA
-0.06065	0.57006	1.36377	-3.47513
ChestPainType_NAP	ChestPainType_TA		
-1.63782	-1.07272		

Degrees of Freedom: 183 Total (i.e. Null); 178 Residual

Null Deviance: 255

Residual Deviance: 172.7 AIC: 184.7

With resampling

```
repeatnew_heart_cv_folds <- vfold_cv(new_heart_train, v= 10, repeats = 5 )  
LR1_resampling <- LR1_wf |>  
  fit_resamples(repeatnew_heart_cv_folds, metrics = metric_set(accuracy, mn_log_loss))  
LR2_resampling <- LR2_wf |>  
  fit_resamples(repeatnew_heart_cv_folds, metrics = metric_set(accuracy, mn_log_loss))  
  
LR1_resampling
```

# Resampling results

# 10-fold cross-validation repeated 5 times

# A tibble: 50 x 5

splits	id	id2	.metrics	.notes
<list>	<chr>	<chr>	<list>	<list>

```

1 <split [660/74]> Repeat1 Fold01 <tibble [2 x 4]> <tibble [0 x 3]>
2 <split [660/74]> Repeat1 Fold02 <tibble [2 x 4]> <tibble [0 x 3]>
3 <split [660/74]> Repeat1 Fold03 <tibble [2 x 4]> <tibble [0 x 3]>
4 <split [660/74]> Repeat1 Fold04 <tibble [2 x 4]> <tibble [0 x 3]>
5 <split [661/73]> Repeat1 Fold05 <tibble [2 x 4]> <tibble [0 x 3]>
6 <split [661/73]> Repeat1 Fold06 <tibble [2 x 4]> <tibble [0 x 3]>
7 <split [661/73]> Repeat1 Fold07 <tibble [2 x 4]> <tibble [0 x 3]>
8 <split [661/73]> Repeat1 Fold08 <tibble [2 x 4]> <tibble [0 x 3]>
9 <split [661/73]> Repeat1 Fold09 <tibble [2 x 4]> <tibble [0 x 3]>
10 <split [661/73]> Repeat1 Fold10 <tibble [2 x 4]> <tibble [0 x 3]>
# i 40 more rows

```

## LR2\_resampling

```

# Resampling results
# 10-fold cross-validation repeated 5 times
# A tibble: 50 x 5
  splits          id    id2    .metrics    .notes
  <list>        <chr>  <chr>  <list>    <list>
1 <split [660/74]> Repeat1 Fold01 <tibble [2 x 4]> <tibble [0 x 3]>
2 <split [660/74]> Repeat1 Fold02 <tibble [2 x 4]> <tibble [0 x 3]>
3 <split [660/74]> Repeat1 Fold03 <tibble [2 x 4]> <tibble [0 x 3]>
4 <split [660/74]> Repeat1 Fold04 <tibble [2 x 4]> <tibble [0 x 3]>
5 <split [661/73]> Repeat1 Fold05 <tibble [2 x 4]> <tibble [0 x 3]>
6 <split [661/73]> Repeat1 Fold06 <tibble [2 x 4]> <tibble [0 x 3]>
7 <split [661/73]> Repeat1 Fold07 <tibble [2 x 4]> <tibble [0 x 3]>
8 <split [661/73]> Repeat1 Fold08 <tibble [2 x 4]> <tibble [0 x 3]>
9 <split [661/73]> Repeat1 Fold09 <tibble [2 x 4]> <tibble [0 x 3]>
10 <split [661/73]> Repeat1 Fold10 <tibble [2 x 4]> <tibble [0 x 3]>
# i 40 more rows

```

## collect\_metrics(LR1\_resampling)

```

# A tibble: 2 x 6
  .metric    .estimator  mean    n std_err .config
  <chr>      <chr>      <dbl> <int>  <dbl> <chr>
1 accuracy  binary      0.675   50 0.00647 Preprocessor1_Model1
2 mn_log_loss binary      0.601   50 0.00571 Preprocessor1_Model1

```

```
collect_metrics(LR2_resampling)
```

```
# A tibble: 2 x 6
  .metric      .estimator mean      n std_err .config
  <chr>        <chr>    <dbl> <int>   <dbl> <chr>
1 accuracy    binary     0.777   50 0.00660 Preprocessor1_Model1
2 mn_log_loss binary     0.488   50 0.00959 Preprocessor1_Model1
```

```
model1 <- collect_metrics(LR1_resampling) |>
  mutate(Model = "Model1")
model2 <- collect_metrics(LR2_resampling) |>
  mutate(Model = "Model2")
combined_metric <- rbind(model1, model2)

combined_metric
```

```
# A tibble: 4 x 7
  .metric      .estimator mean      n std_err .config      Model
  <chr>        <chr>    <dbl> <int>   <dbl> <chr>      <chr>
1 accuracy    binary     0.675   50 0.00647 Preprocessor1_Model1 Model1
2 mn_log_loss binary     0.601   50 0.00571 Preprocessor1_Model1 Model1
3 accuracy    binary     0.777   50 0.00660 Preprocessor1_Model1 Model2
4 mn_log_loss binary     0.488   50 0.00959 Preprocessor1_Model1 Model2
```

I think my best performing model is Model 2, because it has the higher accuracy (.773) and the lowest log loss (.489) means.

## 2 Confusion Matrix

```
LR2_fit <- fit(LR2_wf, data = new_heart_train)
prediction <- predict(LR2_fit, new_data = new_heart_test) |>
  pull() |>
  factor(levels = c("No", "Yes"))
confusionMatrix(
  data = prediction,
  reference = new_heart_test$heartdisease_yes_no
)
```

## Confusion Matrix and Statistics

```

      Reference
Prediction No Yes
      No   73   16
      Yes  21   74

      Accuracy : 0.7989
      95% CI : (0.7336, 0.8543)
      No Information Rate : 0.5109
      P-Value [Acc > NIR] : 5.393e-16

      Kappa : 0.5981

      Mcnemar's Test P-Value : 0.5108

      Sensitivity : 0.7766
      Specificity : 0.8222
      Pos Pred Value : 0.8202
      Neg Pred Value : 0.7789
      Prevalence : 0.5109
      Detection Rate : 0.3967
      Detection Prevalence : 0.4837
      Balanced Accuracy : 0.7994

      'Positive' Class : No
```

The accuracy is about 80%.

### 3 Sensitivity and Specificity and interpretation

Sensitivity : 0.7766

Specificity : 0.8222

Sensitivity, 77.66% is the percent of people who don't have heart disease that receive the negative result. Specificity, 82.22% is the percent of people who truly have heart disease that are correctly identified. This reminds me of false positives and false negatives in inference.